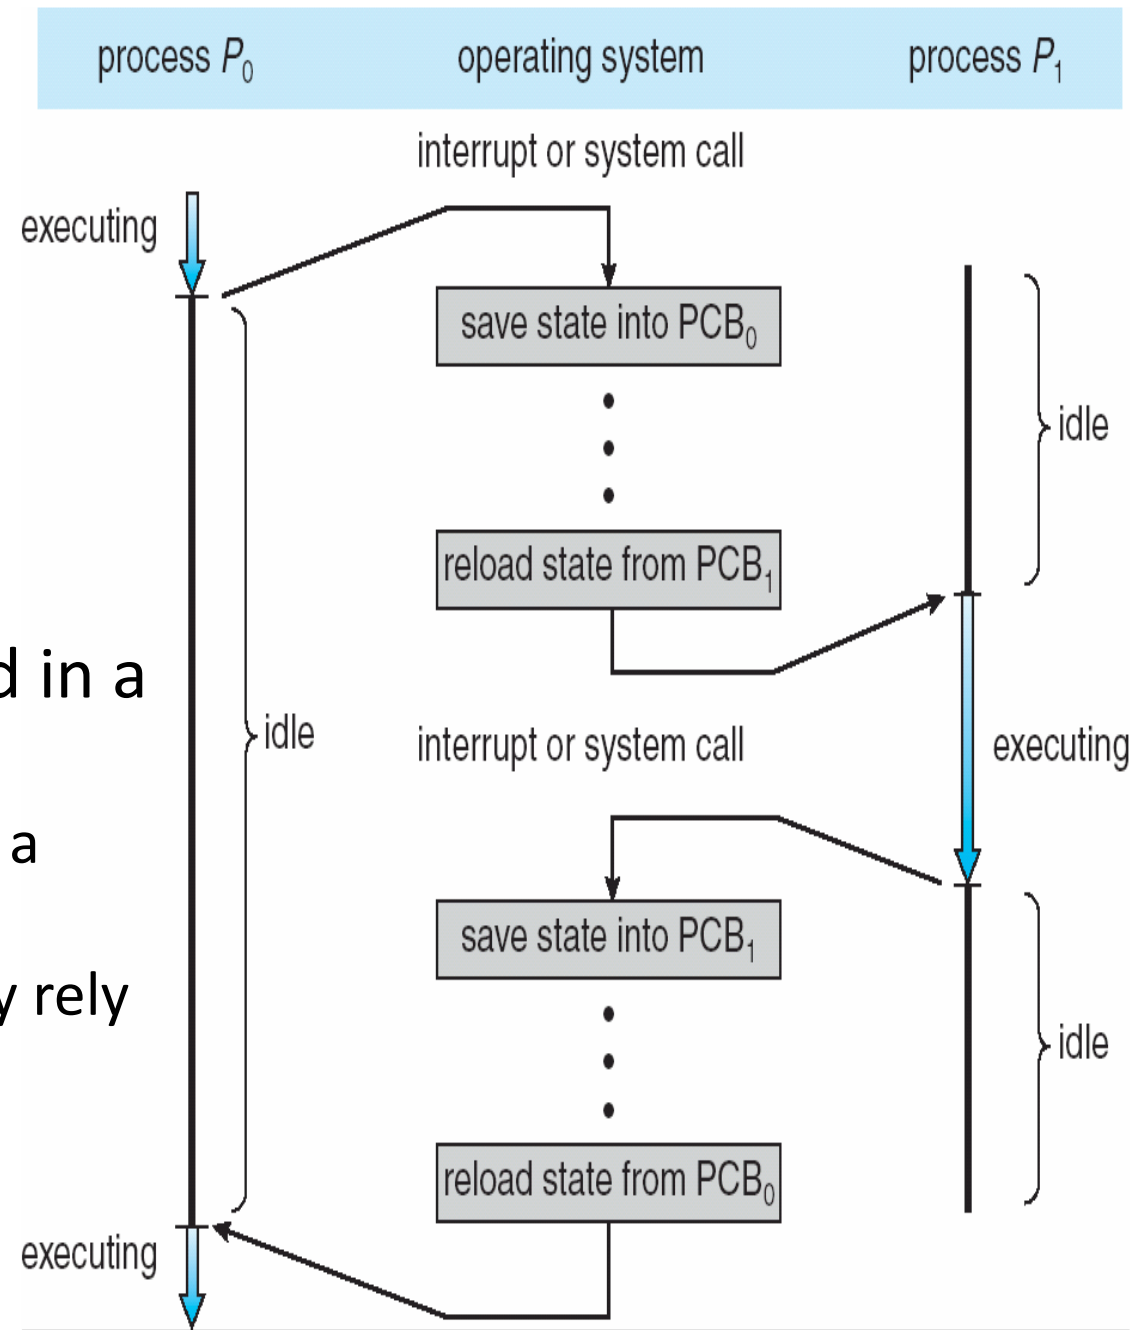


Agenda

- Assignment 1 is out, due September 29!
- Today's lecture
 - Role of the Scheduler and Scheduling Criteria
 - CPU Scheduling Algorithms
 - FCFS
 - Shortest Job First
 - RR
 - Priority Scheduling
 - Multilevel queues
- Reading: Sections 5.1-5.3
- Next Week: Threads and Multiprogramming chapters 3,4

Review

- Not all processes can run at the same time. Why?
- Processes run based on a schedule
- Those that are ready to run are stored in a ***Ready queue***
 - Abstract data structure implemented as a Queue, priority queue, etc...
 - Scheduler and scheduling algorithm may rely on more than one queue



Queues

- Used to schedule processes
- While a process is waiting in the queue, it is in the waiting state
 - **Job** queue – set of all processes
 - **Ready** queue – processes in main memory, ready to execute
 - **Device** queues – processes waiting for an I/O device
- Processes move among various queue creating a cycle

The CPU- I/O Burst Cycle

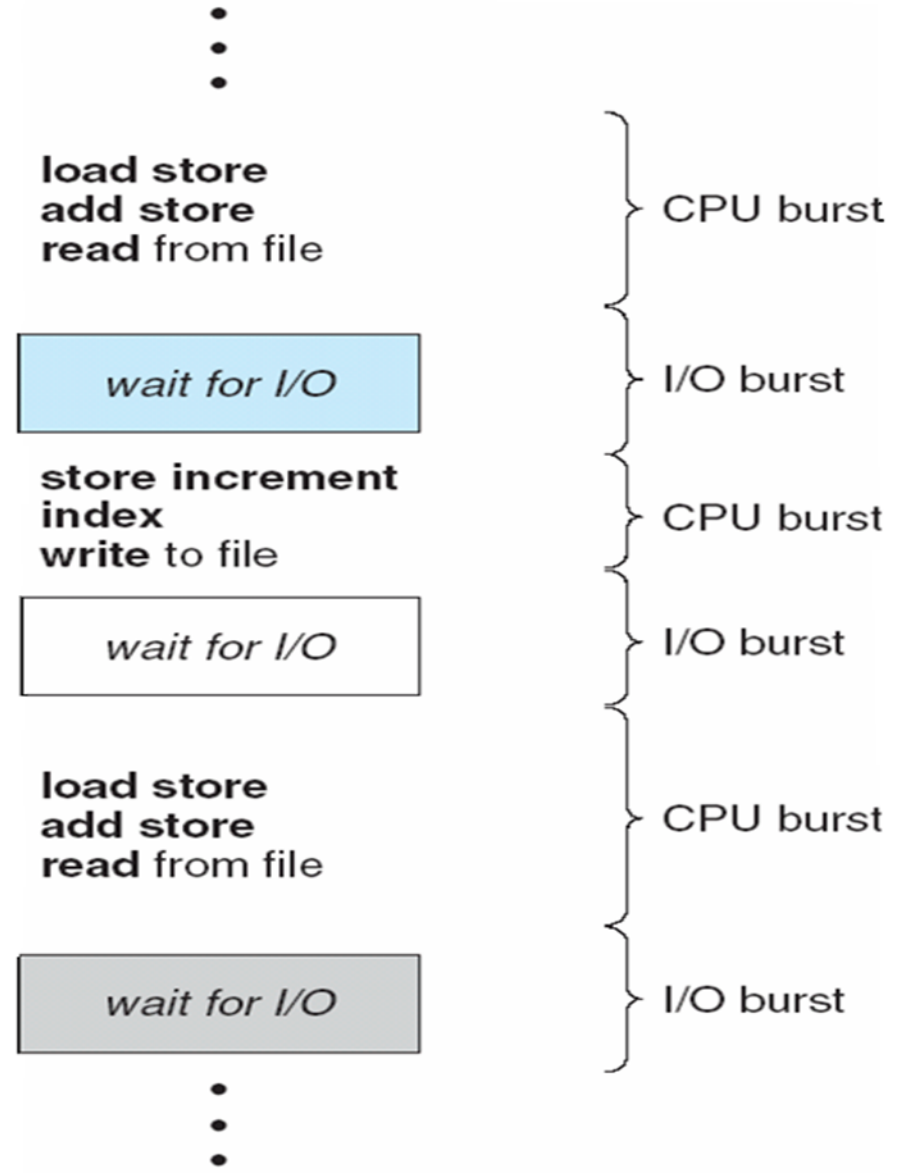
Most processes have a very common execution pattern

CPU burst : perform computation on the CPU

- Execute instructions
- Access memory
- Does not require OS services


I/O burst : performs an I/O operation

- Access a file, network, or other device
- Make a system call to kernel
- Kernel initiates request to hardware
- Hardware is typically slow compared to CPU
- Process waits until hardware request completes
- Process is blocked until request completes

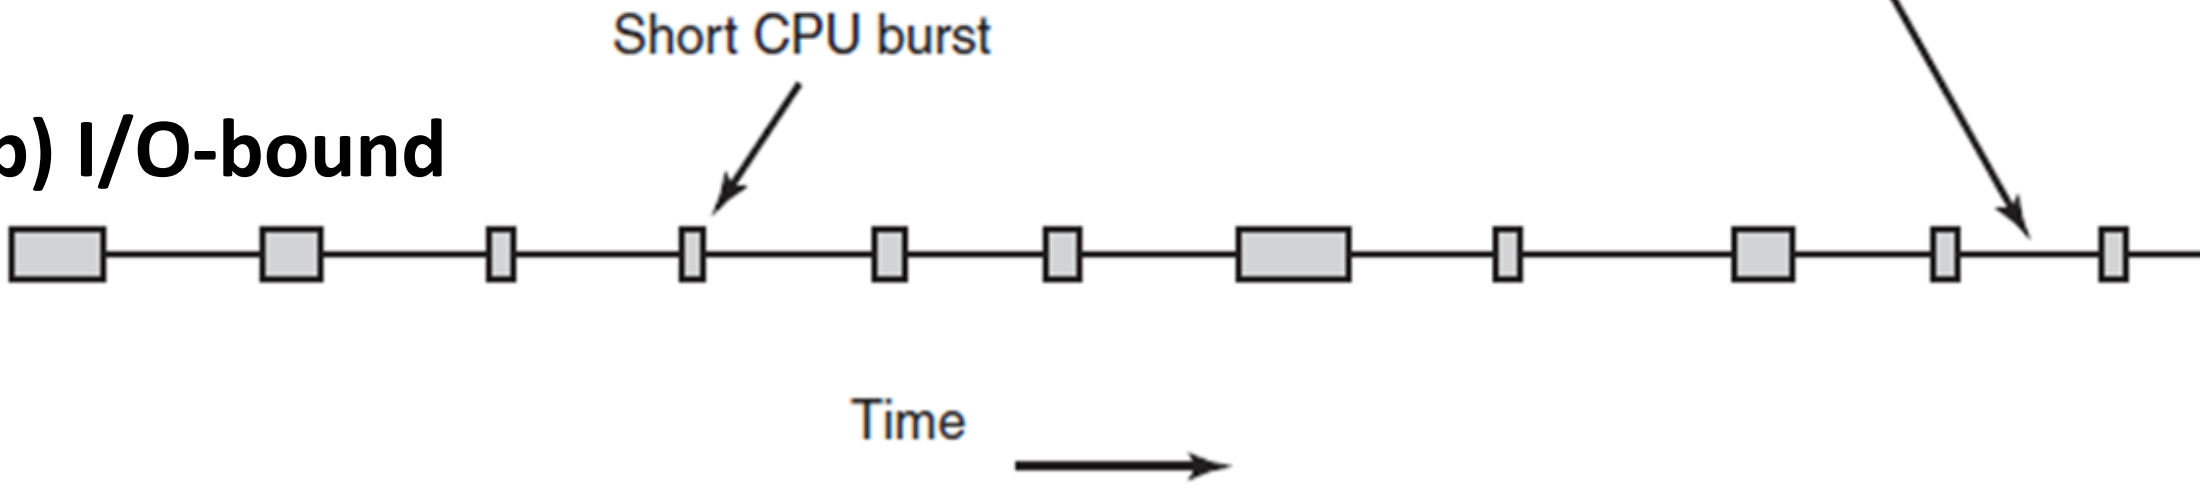


Process Behavior

a) CPU-bound



b) I/O-bound



CPU Scheduling

- During an I/O burst, a process is blocked and is not using the CPU
- We want to maximize the use of the CPU and let it run another process
- Questions:
 - Which process do we choose? There may be many ready processes
 - What happens when waiting process becomes unblocked
 - What if a process does not perform any I/O?
- These issues are addressed by the scheduler

Scheduling Criteria


- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state (I/O or timer interrupt)
 3. Switches from waiting to ready
 4. Terminates
- The type of scheduling algorithm determines when a scheduler runs:
 - Scheduling under 1 and 4 is **non-preemptive (cooperative)**
 - All other scheduling is **preemptive**

Non-preemptive Scheduling

- Scheduling decisions occur when:
 - A running process blocks
 - A running process terminates
- Processes are permitted to run until they request a kernel service
E.g., perform I/O
- Characteristics
 - Efficient (no unnecessary overhead)
 - Simple to implement (simple kernel structure)
 - Predictable, works well for real-time systems
 - Allows processes to monopolize CPU (happened on older OSes)



Preemptive Scheduling

- Processes may be switched (preempted)
 - I/O interrupt
 - Completion of I/O operation
 - Process moves from waiting to ready queue
 - Process moves from waiting to running
 - Timer interrupt
 - Preprogrammed times (10-100 ms)
 - Quantum: the max amount of time that a process can remain on the CPU
- Characteristics
 - Timer servicing adds extra overhead
 - Harder to implement, any interrupt can result in a process switch
 - Less predictable
 - Safer in multiuser systems



...for interrupting...

First-Come, First-Served (FCFS)

- Suppose that the processes arrive in the order: P_1, P_2, P_3
- The Gantt Chart for the schedule is:




Process	Burst Time
P1	24
P2	3
P3	3

- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

First-Come, First-Served (FCFS)

- Suppose that the processes arrive in the order: P_2, P_3, P_1
- The Gantt Chart for the schedule is:



Process	Burst Time
P1	24
P2	3
P3	3

- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect short process behind long process

FCFS


- Pros:
 - Simple FIFO queue needed for implementation
 - Processes at head of queue gets to run next
 - Processes continue to run until it requests a service from the kernel
- Cons:
 - Average waiting time can be long
 - Non-pre-emptive, poor for multi-user systems
- But, can be enhanced when used with pre-emption (Round Robin)

Shortest-Job-First (SJF)

- Associate with each process the length of its next CPU burst.
- Use these lengths to schedule the process with the shortest time
- Pros:
 - Simple priority queue needed for implementation
 - Processes at head of queue gets to run next
 - Process' priority is based on its CPU burst length
 - Processes continue to run until it requests a service from the kernel
 - Optimal – gives minimum average waiting time for a given set of processes
- Cons:
 - Hard to implement
 - Non-pre-emptive, poor for multi-user systems
 - The difficulty is knowing the length of the next CPU request

SJF Example

- SJF scheduling chart



- Average waiting time = $(3 + 16 + 9 + 0)/4 = 7$

Process	Burst Time
P1	6
P2	8
P3	7
P4	3

Determining Length of the Next CPU Burst

- Problem: don't know the length of the process' next CPU burst
- Can only estimate the length based on history
- Can be done by using the length of previous CPU bursts, using exponential averaging

Exponential Moving Average

Use: $\tau_{n+1} = \alpha t_n + (1-\alpha)\tau_n$

- τ_{n+1} = predicted value for the next CPU burst
 - τ_n = Where $0 < \tau_n \leq$ actual length of the nth CPU burst
 - $\alpha < 1$
 - The larger the α the more weight the most recent term has.
-
- Note: For a regular average $\alpha = n^{-1}$
 - Why is this called the exponential moving average?
 - If we expand τ_n
$$\tau_n = \alpha t_n + (1-\alpha)\alpha t_{n-1} + (1-\alpha)^2 \alpha t_{n-2} \dots = \sum_{i=0}^n (1-\alpha)^i \alpha t_{n-i}$$
 - Each previous term is weighed exponentially less than the next
 - Typical value of $\alpha = 1/2$

Exponential Averaging Examples



- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:
$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots$$
$$+ (1 - \alpha)^j \alpha t_{n-j} + \dots$$
$$+ (1 - \alpha)^{n+1} \tau_0$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process

Implementing Priority Scheduling

- Problem: Standard Priority Queue are expensive
 - $O(\log n)$ access time
 - Tree based or heap based
 - Inserting / removing is “complex”
- Idea: Use multiple queues
 - Priorities are fixed -10... 10
 - Use one queue per priority
 - Check each queue for next process
 - E.g. P4 is the next process to run
- Aging
 - Process at head of each queue is promoted to next queue each quantum

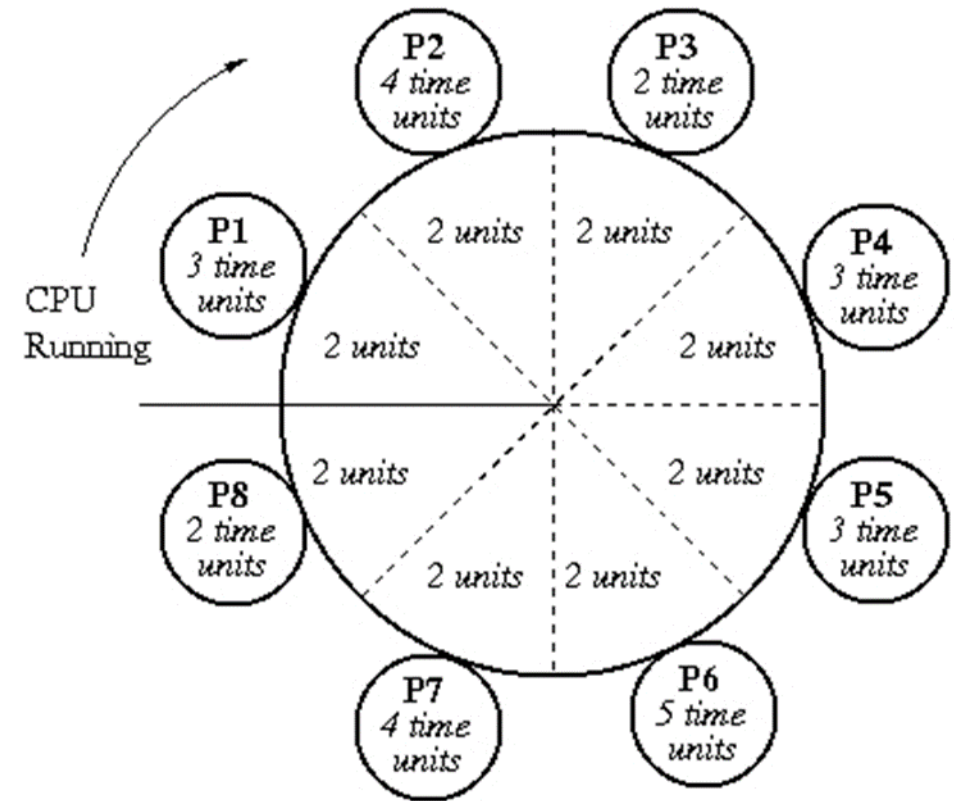


Round Robin Scheduling

- Basic Idea:
 - Single FIFO ready queue of all processes
 - Process at head of queue gets to run next
 - Process continues to run until it
 - Blocks on I/O
 - Terminates
 - Uses up its quantum of execution (timer interrupt goes off)
- Characteristics
 - Simple implementation
 - Indiscriminate between I/O bound and CPU bound processes
 - High waiting time
 - Convoy effect


Round Robin Scheduling

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.



RR Example

- The Gantt chart is:



Process	Burst Time
P1	24
P2	3
P3	3


- Typically, higher average turnaround than SJF, but better response

Time Quantum and Context

Performance

q large \Rightarrow FIFO

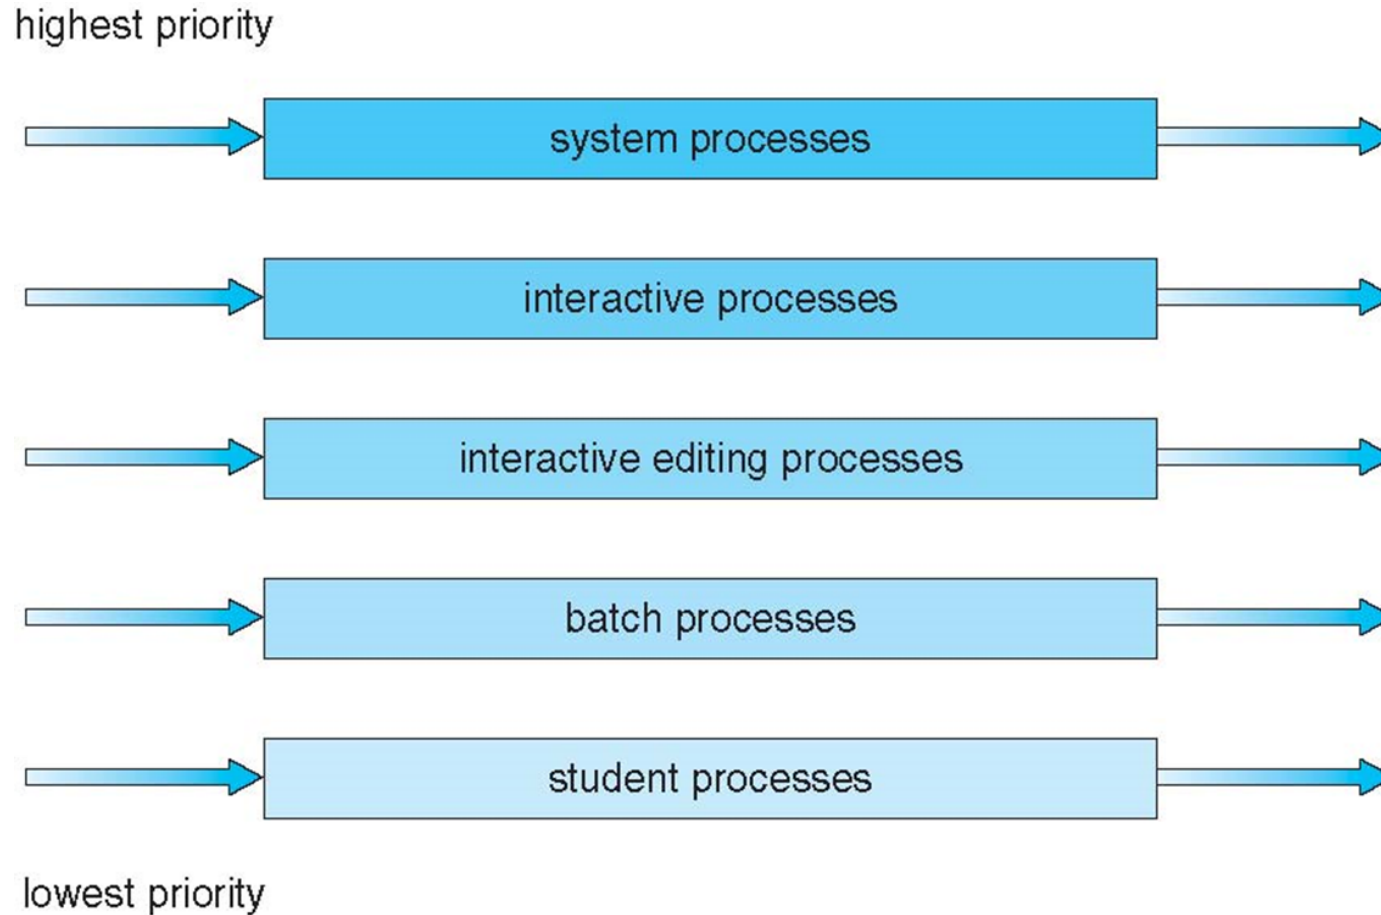
q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high



Multilevel Queue Scheduling

- Ready queue is partitioned into separate queues:
 - Foreground (interactive)
 - Background (batch)
- Each queue has its own scheduling algorithm
 - Foreground – RR
 - Background – FCFS
- Scheduling must be done between the queues
 - Fixed priority scheduling → Possibility of **starvation**
 - Time slice - each queue gets a certain amount of CPU time which it can schedule amongst its processes

Multilevel Queue Scheduling



Multilevel Queue Scheduling

- Characteristics
 - All process get to run
 - More important processes can be run more often
 - If we assign processes correctly, we get low waiting times
 - Used in many systems today
- **Problem:** priorities must be determined a priori
- System cannot adjust to what the process does

Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service


Multilevel Feedback Queue Example

- Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS

- Scheduling

- A new job enters queue Q_0 which is served RR. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
- At Q_1 job is again served RR and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .



What's Next?

- Now that we have a better idea on how the OS schedules processes, we have more questions to address:
- What to do with these processes?
- Why is having multiple processes useful?
- How to make sure processes work together without interfering with each other's resources?